

A hybrid approach to extract protein–protein interactions

Quoc-Chinh Bui^{1,*}, Sophia Katrenko² and Peter M. A. Sloot¹¹Computational Science and ²Theory of Computer Science, Informatics Institute, University of Amsterdam, Science Park 904, Amsterdam, The Netherlands

Associate Editor: Jonathan Wren

ABSTRACT

Motivation: Protein–protein interactions (PPIs) play an important role in understanding biological processes. Although recent research in text mining has achieved a significant progress in automatic PPI extraction from literature, performance of existing systems still needs to be improved.

Results: In this study, we propose a novel algorithm for extracting PPIs from literature which consists of two phases. First, we automatically categorize the data into subsets based on its semantic properties and extract candidate PPI pairs from these subsets. Second, we apply support vector machines (SVMs) to classify candidate PPI pairs using features specific for each subset. We obtain promising results on five benchmark datasets: AIMed, BioInfer, HPRD50, IEPA and LLL with *F*-scores ranging from 60% to 84%, which are comparable with the state-of-the-art PPI extraction systems. Furthermore, our system achieves the best performance on cross-corpora evaluation and comparative performance in terms of computational efficiency.

Availability: The source code and scripts used in this article are available for academic use at <http://staff.science.uva.nl/~bui/PPIs.zip>

Contact: bqchinh@gmail.com

Received on August 11, 2010; revised on October 8, 2010; accepted on October 31, 2010

1 INTRODUCTION

Extraction of protein–protein interactions (PPIs) from literature is an important research topic in the field of biomedical natural language processing (NLP; Miwa *et al.*, 2010). Numerous PPIs have been manually curated and stored into databases, such as BIND, HPRD and IntAct. However, this task has been proven time and resource consuming. As a consequence, most data on PPIs can only be found in literature (Cusick *et al.*, 2009).

Several approaches for extracting PPIs from biomedical text have been reported. These methods range from co-occurrence to more sophisticated machine learning (ML) systems augmented by NLP techniques. Co-occurrence is the simplest approach, which results in high recall but low precision. Rule- or pattern-based approaches can increase precision but significantly lower recall. In addition, these rule sets are derived from training data and are therefore not always applicable to other data they are not developed for (Ananiadou *et al.*, 2006; Kabiljo *et al.*, 2009).

Recently, many ML-based methods have employed NLP techniques such as shallow parsing or full parsing (Björne *et al.*, 2010; Giles and Wren 2008; Giuliano *et al.*, 2006; Miwa *et al.*, 2010). Since full parsing produces more elaborate syntactic information than shallow parsing, PPI extraction systems based on full parsing can potentially yield better results (Miyao *et al.*, 2009). The output of the parser can be represented either as constituent trees or dependency trees. In this case, the PPI extraction task is treated as a binary classification problem which requires a formal protein pair representation and a suitable ML method. A protein pair (an instance) can be represented by a set of features, which are derived from the sentence or its syntactic structure. A ML method is then used to distinguish between positive and negative instances (Niu *et al.*, 2010; Sætre *et al.*, 2010).

Many linguistic features and ML methods have been proposed for the PPI extraction task. Based on feature types, these approaches can be divided into three groups. The first group focuses on lexical and word context features. Bunescu and Mooney (2005) designed a subsequence kernel which uses the following information in a sentence: before the first protein, between two proteins and after the second protein, and combined these features to obtain patterns. Giuliano *et al.* (2006) extended this approach by using a bag-of-words (BOWs) and adding a local context kernel. The second group exploits syntactic features of a sentence. Sætre *et al.* (2007) used various syntactic path features and context features related to words before, between and after two interacting entities. Katrenko and Adriaans (2007) proposed a method based on information found in the predefined levels of the dependency trees, such as local dependency contexts of the protein names and tree's roots. Kim *et al.* (2008) enhanced previous work by proposing a walk kernel which explores the shortest dependency path between two proteins and a modified dependency tree with the parts-of-speech (POS) features. As an alternative to previous approaches, Airola *et al.* (2008) introduced an all-paths graph kernel. They represented a sentence with a dependency graph and considered dependencies connecting two entities outside the shortest path as well as on the shortest path.

Along with the proposed methods, Fayruzov *et al.* (2009), Niu *et al.* (2010) and Van Landeghem *et al.* (2008) also studied individual impact of a variety of feature types on the PPI extraction task. In addition, a study of Miyao *et al.* (2009) has shown that the accuracy of syntactic parsers also contributes to the overall performance of the PPI systems. To compensate for the limitations of each individual feature set and parser errors, Miwa *et al.* (2009a, b) proposed a method that combines all the lexical and parsing features using multiple kernels and parsers. Their system achieved the state-of-the-art performance on a number of benchmark datasets. However,

*To whom correspondence should be addressed.

the studies by Fayruzov *et al.* (2009) and Kim *et al.* (2010) also demonstrated that if two feature types have overlapping rather than complimentary effects, dropping one of them can result in a computationally more efficient method and potentially make a mining algorithm more robust. This argument was also confirmed in the work of Miwa *et al.* (2009b) who showed that excluding the BOWs feature leads to better performance on the AIMed corpus.

Although many approaches have been proposed in the past, the problem of finding the most suitable features for extracting PPIs remain. Adding more features might sometimes improve performance, but they can introduce noise in other cases, or require more computational resources (Kim *et al.*, 2010). In this study, we propose a novel method that consists of two phases. First, we apply semantic rules to partition the dataset into subsets according to its semantic properties and extract candidate PPI pairs from these subsets. Second, we introduce enhanced feature sets for use with a ML classifier to classify these extracted PPI pairs.

To the best of our knowledge, this is the first method that categorizes data into subsets and selects the most appropriate features for each subset. As a result, we increase the robustness of the learning method and make it computationally effective. In general, a PPI extraction system consists of two subtasks: recognizing protein names (NER) and extracting PPI pairs. However, this study only focuses on the PPI extraction task with an assumption that relevant named entities were given.

2 METHODS

The workflow of the proposed system is as follows:

- (1) Text preprocessing.
- (2) Extracting candidate PPI pairs.
- (3) Classifying extracted PPI pairs.

2.1 Text preprocessing

Text preprocessing includes converting protein names using a predefined rule set, filtering out input sentences with only one protein, splitting input sentences contain multiple clauses, and parsing sentences using the Stanford lexical parser (Klein and Manning, 2003).

2.1.1 Processing protein names In order to improve accuracy of the parser, we replace all mentioned protein names with a place holder, i.e. PRO1, PRO2 (we refer to them as PRO*). We define a rule set to resolve the problem with embedded protein names (e.g. AIMed corpus), protein names that share prefix or suffix (e.g. AIMed and BioInfer corpora) and protein names including multiple positions. After this process, the number of proteins in the sentence is not changed. The list of original protein names for each sentence is maintained for reference purposes.

2.1.2 Replacing parenthetical remarks and splitting a sentence If no word inside parentheses is a protein name, then all words and the parentheses are removed. In case the sentence consists of multiple clauses, the system splits it into clauses. The resulting sentence is referred to as a *simplified sentence*.

In the last step of text preprocessing, a simplified sentence that contains at least two protein names is analyzed with the Stanford lexical parser to produce a syntactic tree. All parse trees are stored in a local database for later use. An example of a sentence and its output after text-preprocessing step is given below:

AIMed.d101.s852: *The bZIP domains of Fos and Jun mediate a physical association with the TATA box-binding protein.*

Its simplified sentence: *The bZIP domains of PRO0 and PRO1 mediate a physical association with the PRO2.*

2.2 Extracting candidate PPI pairs

Previous studies (Bui *et al.*, 2010; Fundel *et al.*, 2007; Rinaldi *et al.*, 2010) have shown that, in biomedical text, relation between two entities (protein–protein, protein–gene, drug–mutation and others) can be expressed in the following abstract forms:

Form 1: PRO_i word* REL (verb) word* PRO_j

Example: *PRO1 interacts with PRO2.*

Form 2: PRO_i word* REL (noun/verb) word* PRO_j

Example: *PRO1 has a weak interaction with PRO2.*

Form 3: REL (noun) word* PRO_i word* PRO_j

Example: *interaction between PRO1 and PRO2.*

Here, REL is a cue word (interaction, inhibit, etc.) and can be a noun or a verb, word* are tokens between PRO* and REL. PRO_i and PRO_j can be any protein pair with $j \geq i + 1$ (e.g. <PRO1, PRO2>). In addition, a closer look at the annotated corpora (e.g. AIMed and BioInfer) reveals that we can define two more forms:

Form 4 (*compound form*): PRO_i/PRO_{i+1} or PRO_i PRO_{i+1} or PRO_i-PRO_{i+1}, if two entities appear in the sentence with the patterns above, they seem to have interaction.

Example: *PRO1/PRO2 binding; PRO1-PRO2 compound.*

Form 5 (*complex form*): PRO_i word* PRO_j word* REL

Example: *PRO1, PRO2 interact; in PRO1, PRO2 complex.*

Form 4 is expressed as a pattern and requires an exact match. For other forms, there might be one or more tokens between <PRO_i, REL>, <REL, PRO_j> or <PRO_i, PRO_j>. Based on these basic forms, we now map the semantic relations of these forms into parse trees. For convenience, we define the following patterns:

- (a) *Full clause*: S < ((NP << PRO*) \$++ (VP << PRO*))
- (b) *Partial clause*: VP < ((NP << PRO*) \$++ (S < (VP << PRO*)))
- (c) *Subclause*: S < ((NP << PRO*) \$++ (VP < (S|SBAR << PRO*)))
- (d) *NP pattern*: NP < ((NP << # REL) \$++ (PP << PRO*))

Here S (clause), NP (noun phrase), VP (verb phrase), PP (prepositional phrase) and SBAR (subclause) are constituents of the parse tree, PRO* is a place holder for a protein name, and REL is a cue word of the input sentence. All patterns above are written using the Tregex syntax (<http://nlp.stanford.edu/software/tregex.shtml>), which is developed within the Stanford parser package. Figure 1 illustrates some parse trees with the patterns mentioned above, e.g. the parse tree in Figure 1a shows a *full clause*, Figure 1b shows a *NP pattern* and Figure 1c shows a *subclause*.

2.2.1 Relation list We created a relation list by combining the relation lists used in the previous work by Chowdhary *et al.* (2009), Fundel *et al.* (2007) and Niu *et al.* (2010).

In the following section, we describe the procedures and algorithm to extract candidate PPI pairs from a parse tree or a simplified sentence:

Procedure for Form 1: this procedure requires a parse tree which contains *full clause*, *partial clause* or *subclause* patterns. The procedure is as follows:

- (a) Check a head word of VP that corresponds to satisfied pattern. If this verb belongs to the relation list, use it as REL.
- (b) Create a *right_keys* list: use all keys in VP as *right_keys*.

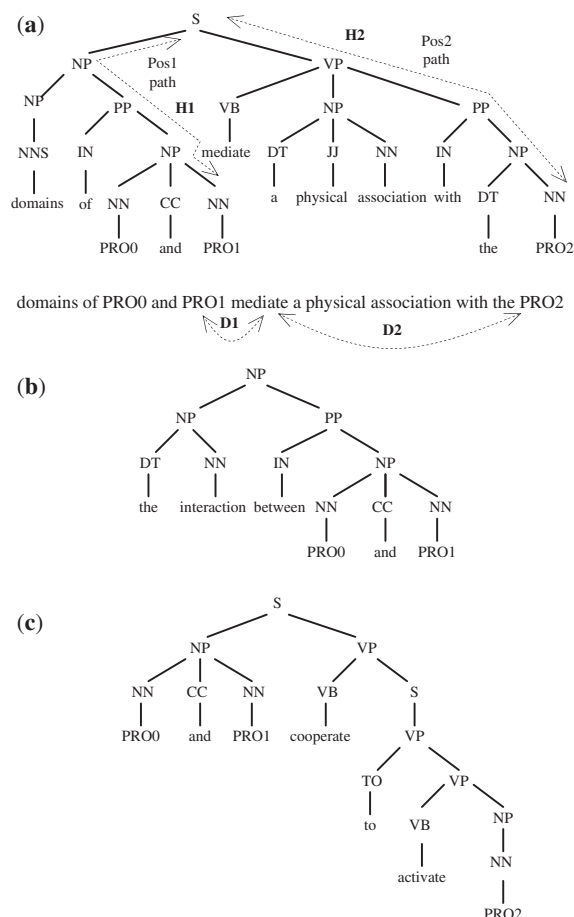


Fig. 1. Examples of parse trees output from the Stanford lexical parser: (a) parse tree satisfies full clause pattern, (b) parse tree satisfies NP pattern and (c) parse tree satisfies subclause pattern. Some features (D1, D2, H1, H2) and paths from the join node connecting a given protein pair (PRO1, PRO2) are also shown in (a).

(c) Create a *left_keys* list: use all keys in NP if a satisfied pattern is *full clause* or *subclause*. If a pattern is *partial clause*, then find a sister NP immediately preceding VP.

(d) Form candidate PPI pairs: enumerate the *left_keys* and *right_keys* to compose a triple $\langle \text{PRO}_i\text{-REL-PRO}_j \rangle$.

Procedure for Form 2: form candidate PPI pairs from a simplified sentence if they satisfy the following form: $\text{PRO}_i \text{ word* REL word* } \text{PRO}_j$.

Procedure for Form 3: this procedure requires a parse tree which contains NP pattern. The procedure is as follows:

- Check a head word of NP that corresponds to satisfied pattern. If this noun belongs to the relation list, use it as REL.
- Find a proposition as a splitter in pattern's PP phrase in order to create *left_keys* and *right_keys* lists.
- Form candidate PPI pairs: enumerate the *left_keys* and *right_keys* to compose a triple $\langle \text{PRO}_i\text{-REL-PRO}_j \rangle$.

Procedure for Form 4: form candidate PPI pairs from a simplified sentence if they satisfy any of the following forms: $\text{PRO}_i/\text{PRO}_{i+1}$; $\text{PRO}_i \text{ PRO}_{i+1}$; $\text{PRO}_i\text{-PRO}_{i+1}$.

Procedure for Form 5: form candidate PPI pairs from a simplified sentence if they satisfy the following form: $\text{PRO}_i \text{ word* } \text{PRO}_j \text{ word* REL}$, and if the distance from the PRO_i to the REL is shorter than five tokens.

For these procedures, the extracted candidate PPI pairs obtained from the same procedure are grouped together, i.e. outputs from Form 1 are grouped into group 1. Among these groups, group 2 overlaps group 1 when REL is a verb, with the purpose to recover PPI pairs that the group 1 failed to extract due to the parser errors (Miyao *et al.*, 2009). As a consequence, in some cases, a PPI pair may belong to more than one group. Therefore, the order in which patterns are applied is important (as described in the Algorithm 1 below).

Algorithm 1. // Algorithm to extract candidate PPI pairs.

Input: simplified sentence, parse tree and the relation list

Output: list of candidate PPI pairs of corresponding groups

Init: *used_list* = null // store extracted pairs to avoid overlap since one pair can satisfy more than one pattern.

For a list of subsentence/parse tree

For form from 1 to 5

Extract candidate pairs from parse tree or simplified sentence

If candidate pairs found

For each candidate pair

Check whether this pair in the *used_list*

If not found

Store this pair to the corresponding group

Store this pair to the *used_list*

End if

End for

End if

End for

End for

2.3 Features

In this section, we propose feature sets for ML classification. Our feature sets are combinations of some features that were previously proposed by Chowdhary *et al.* (2009), Giuliano *et al.* (2006), Miwa *et al.* (2009a, b) and Niu *et al.* (2010). For each candidate PPI pair extracted from an input sentence, the system outputs a triple, e.g. $\langle \text{PRO}_i, \text{REL}, \text{PRO}_j \rangle$ then the following features are generated:

Keyword: is a relation word (REL) from the extracted triple. In addition, we also count the number of protein names (C1) and relation words (C2) in each simplified sentence.

Distance: we use two features: D1, D2 (Fig. 1a) to measure the distance (number of words/tokens) between $\text{PRO}_i\text{-REL}$ and REL-PRO_j (or between REL-PRO_i and $\text{PRO}_i\text{-PRO}_j$ for group 3).

Height: we use two features: H1, H2 (Fig. 1a) to measure the distance from the joint node connecting PRO_i and PRO_j in the parse tree. These features are similar to D1 and D2 except that they measure the number of nodes on the paths from a local root to PRO_i and PRO_j .

POS: we use two lists: Pos1, Pos2 (Fig. 1a) to store POS and syntactic features from the joint node connecting PRO_i and PRO_j in the parse tree, respectively.

Lexical: this feature is a modification of a BOWs. Instead of using all tokens, we only consider tokens that belong to the *relation list* and a list of prepositions: *and, or, by, through, in, of, to* and *between*. If a token is a protein (PRO^*), then its value is replaced with 'KEY'. We use four lists of tokens:

L1: a list of tokens between PRO_i and REL, or between REL and PRO_j for group 3.

L2: a list of tokens between REL and PRO_j , or between PRO_i and PRO_j for group 3.

Table 1. List of features and their usage for each group

Feature	Keyword	Distance	Height	POS	Lexical
Group 1	Rel, C1, C2	D1, D2	H1, H2	Pos1,2	L1, L2
Group 2	Rel	D1, D2	H1, H2		L1, L2
Group 3	Rel, C1, C2	D1, D2	H1, H2	Pos1,2	L1, L2
Group 4	C1, C2				L3, L4
Group 5	Rel C1, C2	D1, D2			L3, L4

L3: a list of tokens before PRO_i .

L4: a list of tokens after PRO_j .

Selection of features for each individual group: the important benefit of partitioning data into subsets is that we can select the most appropriate features for each subset. Let us consider the following two cases: a PPI pair in form 4 (compound form) and a PPI pair in form 1 (full clause). For a PPI pair in a compound form, e.g. PRO1-PRO2, the shortest path features proposed in the previous work become useless because no feature can be extracted from this path. In this case, the BOW features seem to be the most appropriate ones. In contrast, for the PPI pairs in form 1, e.g. PRO1 interacts with PRO2, the shortest path, and the tokens between PRO1 and PRO2, play an important role. Based on the properties of each group of extracted PPI pairs, we manually select features that are potentially suitable for that group. Table 1 shows the list of features corresponding to each group. Furthermore, we also use the *Ranker (attribute selection)* method from the WEKA ML package (Witten and Frank, 2005) to determine the length of the POS and lexical lists, which are limited to maximum six attributes.

ML method: Support vector machines (SVMs) have been widely used in the PPI extraction task, and has shown competitive performance over other learning methods (Kim *et al.*, 2010; Sætre *et al.*, 2010). In this work, we use SVM classifier with a default RBF kernel and tune the complexity parameter C by using *CVParameterSelection* function from the WEKA. All individual features mentioned above are combined into a single feature vector for the classification task. Depending on each group of PPI candidates, the number of features ranges from 14 to 26 as shown in Table 1. For example, for a candidate PPI pair <PRO1, PRO2> extracted from group 1 (as shown in Fig. 1a), the following features are generated:

REL: mediate; D1: 1, D2: 6; C1: 3, C2: 2; H1: 5, H2: 4; L1: null, null, null; L2: association, with, null; Pos1: NP, PP, NP, NN, null, null; Pos2: VB, PP, NP, NN, null, null, true.

For more details on text preprocessing, PPI extraction and feature generation steps, see Supplementary source code provided in this article.

3 RESULTS AND DISCUSSION

3.1 Datasets

We use five corpora (<http://mars.cs.utu.fi/PPICorpora/GraphKernel.html>) which have been converted into a unified format and are provided by Pyysalo *et al.* (2008): AIMed, BioInfer, HPRD50, IEPA and LLL. Table 2 shows the quantitative information of all five corpora.

3.2 Evaluation methods

We perform two types of evaluation, a single corpus test and a cross-corpora test. In the single corpus test, we evaluate the performance of the proposed method by 10-fold abstract-wise cross-validation (CV), and use the one-answer-per-occurrence criterion (Miwa *et al.*, 2009b). We split the corpora as recommended by Airola

Table 2. Statistics on five corpora

Corpus	AIMed	BioInfer	HPRD50	IEPA	LLL
Positive pairs	1000	2534	163	335	164
All pairs	5834	9666	433	817	330

Table 3. Results of PPI extraction algorithm on five corpora

Corpus	AIMed		BioInfer		HPRD50		IEPA		LLL	
	TP	FP	TP	FP	TP	FP	TP	FP	TP	FP
Group 1	500	701	849	648	101	38	170	78	106	13
Group 2	143	112	215	154	14	10	26	26	2	4
Group 3	113	322	330	529	17	23	92	61	32	11
Group 4	22	55	170	73	0	8	4	3	0	2
Group 5	39	192	134	182	6	5	3	6	0	1
Total	817	1382	1698	1586	138	84	295	174	140	31
Recall (%)	81.7		67		84.7		88.1		85.4	
Precision (%)	37.2		51.7		62.2		62.9		81.9	
F-score (%)	51.1		58.4		71.7		73.4		83.6	

et al. (2008) in order to directly compare our results against the performance of other systems. In the cross-corpora test, we conduct two experiments. First, we use one corpus for training and test the system on the four remaining corpora. Second, we use four corpora (ALL) for training and test on the remaining corpus.

Precision, recall and *F*-score are used as evaluation metrics. Let TP denote numbers of true positives, FP denote the number of false positives and FN denote the numbers of false negatives. The measures are defined as follows:

$$\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$$

$$\text{Precision} = \text{TP} / (\text{TP} + \text{FP})$$

$$F\text{-score} = 2 * \text{Recall} * \text{Precision} / (\text{Recall} + \text{Precision})$$

3.3 Performance of PPI extraction algorithm

Table 3 shows the results of the PPI extraction algorithm on five corpora. In order to calculate recall for each corpus, we use the number of original positive pairs from Table 1 (e.g. for AIMed, $\text{TP} + \text{FN} = 1000$, for BioInfer, $\text{TP} + \text{FN} = 2534$). It is worth noting that, in some systems (Airola *et al.*, 2008; Kim *et al.*, 2010), self-interaction pairs are removed prior to the evaluation; therefore, the *F*-score can be higher if we adopt these settings.

Table 3 also presents the properties of each extracted group in each corpus and the differences between corpora. Clearly, group 1 is the most common among all corpora and accounts for >50% of TP pairs of all corpora except for BioInfer. The first three groups are also common for all five corpora and account for >80% of the extracted pairs. However, groups 4 and 5 are more corpus-specific and can be found mostly in AIMed and BioInfer corpora. This information provides more insight for understanding the properties of five corpora studied previously by Pyysalo *et al.* (2008).

Table 4. Results of the 10-fold abstract-wise CV on five corpora

Corpus	Precision	ΔP	Recall	F -score	ΔF
AIMed	55.3	18	68.5	61.2	10
BioInfer	61.7	10	57.5	60.0	2
HPRD50	70.2	8	77.9	73.8	2
IEPA	67.4	5	83.9	74.7	1
LLL	84.1	2	84.1	84.1	1

ΔP and ΔF show the increase in precision and F -score compared with their corresponding values in Table 3. The values are shown in percentage (%).

The results in Table 3 demonstrate that the proposed algorithm is capable of extracting of >80% of positive pairs (recall > 80%) on AIMed, HPRD50, IEPA and LLL corpora. This means that these four corpora share some common patterns or the same annotation policy. However, our method can only extract 67% positive pairs from BioInfer. To find out why our algorithm has lower recall on the BioInfer corpus, we examine sample sentences from BioInfer which the system failed to extract. By analyzing these sentences, we discovered that in some cases, BioInfer has a special annotation policy. Consider, for example, the following sentence:

BioInfer.d436.s0: *Moreover, ectopic expression of PRO0 and PRO1 can stimulate the PRO2 promoter in an PRO3-dependent manner, and this is inhibited by coexpression of the PRO4 (PRO5) PRO6.*

For this input sentence, our system extracts only four TP pairs, but the number of positive PPI pairs provided by the BioInfer corpus is 11 (see sentence *BioInfer.d436.s0* in BioInfer corpus for more detail). Therefore, to increase recall on BioInfer, additional study of this corpus is needed.

Table 3 also shows that the performance of algorithm 1 is comparable with other ML-based systems. Moreover, it outperforms the best rule-based system on all five corpora [see Kabiljo *et al.* (2009) for more details].

3.4 Single corpus evaluation

With the extracted PPI pairs obtained in the previous step, we use a standard SVM classifier to classify them. Comparing with the original data from Table 1, the number of positive and negative instances of the AIMed and BioInfer corpora are quite balanced (Table 3). To calculate the F -score, we accumulated all TP and FP pairs from the 10-fold CV test (reported by WEKA via the confusion matrix). For TP+FN values, similarly to Section 3.3, we use the original positive pairs of each corpus in Table 1. For the IEPA, HPRD50 and LLL corpora, however, we only use the first three groups of extracted pairs because the number of candidate pairs in the groups 4 and 5 is too small for the 10-fold CV.

Table 4 shows the performance of the classifier on five corpora. The results indicate that by using our feature vectors, the classifier can further boost the performance of the overall system. When comparing precision on each corpus before and after applying classification, we can see a significant increase in precision (ΔP values). Among all corpora, the AIMed corpus gains the most increase in precision with ΔP of 14%. Even on a small corpus like LLL, with already very high precision, the final precision still gains

Table 5. Performance comparison between full dataset, filtered dataset and partitioned dataset using 10-fold CV on five corpora

Corpus	AIMed	BioInfer	HPRD50	IEPA	LLL
Full dataset	52.5	58.7	68.6	72	83.6
Filtered dataset	55.2	59.6	76.0	72.9	84.4
Partitioned dataset	61.2	60.0	73.8	74.7	84.1

Full dataset and filtered dataset use the same feature sets. The performance is measured by the F -score (%).

Table 6. Performance comparison with other systems on AIMed

System	Description	F -score (%)
Sætre (2010)	Feature-based, two parsers	64.2
Miwa (2009a)	Multiple kernels, two parsers	60.8
Kim (2010)	Walk-weighted subsequence kernels, one parser	56.6
Airola (2008)	All-paths graph kernel, one parser	56.4
Niu (2010)	Linear kernel, one parser	53.5
BKS	RBF kernel, one parser	61.2

2%. For each corpus, recall decreases after applying classification but the final F -scores increase for all corpora and range from 1% to 10%. Furthermore, Table 4 shows that our system achieves the best performance on three corpora: IEPA, HPRD50 and LLL using 10-fold CV when compared with other results reported so far.

To study the benefit of filtering data (obtained after applying rules on the full dataset) and partitioning data, we conduct the experiment in which the full dataset and the filtered dataset use all features in Table 1. However, in this experiment the REL feature is not available, D1 and D1 features are replaced by D1+D2 and L1 and L2 features are merged. Table 5 summarizes the performance of the system on the full dataset (all PPI pairs), filtered dataset and partitioned dataset. The results from Table 5 show that when evaluating on filtering dataset, the system has better performance on all five corpora compared with full dataset. In addition, the performance further improve on three largest corpora (AIMed, BioInfer and IEPA) when we partition data and select feature specific for each sub-dataset. This clearly shows the benefit of our hybrid approach that combines rule with partition data.

Table 6 illustrates a comparison of our system (BKS) against recent work on the AIMed corpus, which is considered the *de facto* benchmark for the PPI extraction task. However, the comparison may not be straightforward because these PPI systems can use different text preprocessing and learning methods. The results show that our system with its F -score of 61.2% is comparable with the state-of-the-art systems proposed by Miwa *et al.* (2009a) and Sætre *et al.* (2010). In addition, our approach significantly outperforms other methods based on one parser's output.

3.5 Cross-corpora evaluation

In addition to the standard 10-fold CV, recent studies also suggested to test existing approaches to PPI extraction using cross-corpora. This type of evaluation can show whether a system trained on one corpus can perform equally well on other corpora, with an

Table 7. Results of the cross-corpora test on five corpora

Corpus	AIMed		BioInfer		ALL	
	P	F	P	F	P	F
AIMed	–	–	44.4	55.2	44.1	55.1
BioInfer	57.1	54.4	–	–	64.3	50.5
HPRD50	67.0	72.3	69.9	74.0	68.8	70.8
IEPA	67.3	73.6	68.6	73.8	70.1	72.9
LLL	85.7	83.0	85.0	82.4	86.0	80.1

Columns correspond to training and rows correspond to testing corpora. ALL refers to a situation where four corpora are used for training and the remaining corpus for testing. Precision (P) and *F*-score (F) are shown in percentage (%).

assumption that these corpora addressed the same PPI task. In this setting, one corpus is used for training and the remaining corpora are used for testing. For the evaluation, we only use the two largest corpora, i.e. AIMed and BioInfer, for training since the other corpora are too small after partitioning onto five groups. Furthermore, we conduct the experiment proposed by Fayruzov *et al.* (2009) and Kabiljo *et al.* (2009), in which four corpora are used for training and the remaining corpus for testing.

Table 7 shows the results of the cross-corpora evaluation on five corpora. Here, the columns correspond to the training set and the rows correspond to test sets. Table 7 illustrates that our system achieves the best performance when BioInfer is used as the training set. Note that we do not use data from group 4 to group 5 due to their corpus-specific properties; therefore, the performance might be higher if these groups are used. This finding is also consistent with the evaluation results by Miwa *et al.* (2009a) and Airola *et al.* (2008). In their work, better performance is also obtained with BioInfer being used for training.

Interestingly, Table 7 shows that even though the *F*-score on each corpus in the cross-corpora test is lower than *F*-score in the single corpus test, the precision is significantly higher when compared against the initial values in Table 3. This means that performance on all corpora can be boosted when precision is given a priority. This also implies that the classifier is able to learn from cross-corpora training data. In other words, the proposed feature vectors are compatible across the corpora.

Table 8 shows the performance of our system (BKS) in the cross-corpora setting compared with other approaches. The results in Table 8 demonstrate that BKS outperforms others when AIMed and BioInfer corpora are used as training sets. In addition, for all evaluation tests, we use the same settings for training corpora except for the complexity term (*C* parameter of the RBF kernel). This is one step closer to the real world situation. However, Miwa *et al.* (2009b) have obtained higher performance compared with ours by applying corpus weighting.

3.6 Computational performance

Another important aspect in PPI extraction is the computational time taken to train and test the system. A previous study by Van Landeghem *et al.* (2008) has shown that in order to reduce the number of computational resources used by ML, one has to consider a trade-off between performance (in this case measured by the *F*-score) and the computational time required by ML. Despite this

Table 8. Performance comparison with other systems using cross-corpora evaluation

	AIMed			BioInfer		
	BKS	Miwa (2009a)	Airola (2008)	BKS	Miwa (2009a)	Airola (2008)
AIMed	–	–	–	55.2	49.6	47.2
BioInfer	54.4	53.1	47.1	–	–	–
HPRD50	72.3	68.3	69.0	74.0	68.3	63.9
IEPA	73.6	68.1	67.4	73.8	71.4	68.0
LLL	83.0	73.5	74.5	82.4	76.9	78.0

Columns correspond to training and rows correspond to test sets. The performance is measured by the *F*-score (%).

Table 9. Computational time for 10-fold CV of full and partitioned dataset of the AIMed corpus

Dataset	WEKA – RBF	LibSVM – RBF
Full dataset (s)	10 812	194
Partitioned dataset (s)	77	39

These values exclude time for parsing and text preprocessing. Two RBF kernels from WEKA and LibSVM are used. The experiment is run on a PC with an Intel 3.2 GHz processor and 4 GB of RAM.

fact, only few systems report on how many computational resources their systems use for training and testing. Miwa *et al.* (2009b) demonstrate that by using suitable features as well as a learning method, they can improve the performance of their previously proposed system (Miwa *et al.*, 2009a). In addition, Fayruzov *et al.* (2009) have pointed out that the more kernels the system uses the more computational resources are needed.

Table 9 shows the running time of our system on the full dataset compared against partitioned dataset of the AIMed corpus. The system runs much faster on the partitioned dataset with both RBF kernels from WEKA and LibSVM (<http://www.csie.ntu.edu.tw/~cjlin/libsvm/>). In addition, since our method partitions data into five groups, we also test it in a node with eight CPUs (Xeon 3.0 GHz). In this test, we use the RBF kernel from LibSVM and the classifiers are run in parallel. The maximum time used by the system is 12 s. This means that our system not only runs fast on a single PC, but can also be used in parallel, which is particularly suited for large-scale experiments.

4 CONCLUSIONS

In this article, we propose a novel method for extracting PPI pairs from literature. Our approach combines the strength of both semantic rules and ML classification. The evaluation on five benchmark corpora has shown that our system achieves results comparable with the best PPI extraction methods on a single corpus. It outperforms other systems on cross-corpora test and has fast running time.

The proposed method consists of two phases: partitioning data into subsets then extract candidate PPI pairs from these subsets, and classifying extracted PPI pairs. The advantages of this method are

4-fold. First, it filters out a significant amount of negative (non-interaction) PPI pairs, thus balancing the number of positive and negative pairs in training data. Second, by partitioning data into subsets, we can select the most appropriate features for each subset, which potentially improves the final performance. Third, our system only uses a small set of features and therefore performs the best in terms of computational resources. Finally, the classifier can be run in parallel on each subset, which is desirable for the large-scale experiments.

In addition, our method uses five semantic rules; therefore, it is generic and can be easily applied to new datasets. Furthermore, it is easy to set-up because it only uses publicly available NLP tools and a standard ML package. In addition, the proposed method can also be extended to extract new relation types in biomedical text, e.g. complex event extraction (Björne *et al.*, 2010, Miwa *et al.*, 2010, Rinaldi *et al.*, 2010).

For future work, we plan to integrate a NER tagger into our system in order to study the effect of its performance when a gold-standard NER is not given. In addition, the performance of the overall PPI extraction task largely depends on the output of the extraction phase. Increasing recall in this phase would further boost the performance, especially in case of the BioInfer corpus.

ACKNOWLEDGEMENTS

The authors sincerely thank Dr Makoto Miwa and Rick Quax for their useful comments and suggestions.

Funding: This research was supported by the European Union through the DynaNets project <http://www.dynanets.org> EU grant agreement no: 233847; the Vietnamese Oversea Training Program, and the VL-e project (BSIK grant).

Conflict of Interest: none declared.

REFERENCES

- Airola, A. *et al.* (2008) All-paths graph kernel for protein-protein interaction extraction with evaluation of cross-corpus learning. *BMC Bioinformatics*, **9**, S11.
- Ananiadou, S. *et al.* (2006) Text mining and its potential applications in systems biology. *Trends Biotechnol.*, **24**, 571–579.
- Björne, J. *et al.* (2010) Complex event extraction at PubMed scale. *Bioinformatics*, **26**, i382–i390.
- Bui, Q.C. *et al.* (2010) Extracting causal relations on HIV drug resistance from literature. *BMC Bioinformatics*, **11**, 101.
- Bunescu, R. and Mooney, R. (2005) Subsequence kernels for relation extraction. In *Proceedings of the 19th conference on Neural Information Processing Systems*. Vancouver, Canada.
- Chowdhary, R. *et al.* (2009) Bayesian inference of protein-protein interactions from biological literature. *Bioinformatics*, **25**, 1536–1542.
- Cusick, M.E. *et al.* (2009) Literature-curated protein interaction datasets. *Nat. Methods*, **6**, 39–46.
- Fayruzov, T. *et al.* (2009) Linguistic feature analysis for protein interaction extraction. *BMC Bioinformatics*, **10**, 374.
- Fundel, K. *et al.* (2007) RelEx - Relation extraction using dependency parse trees. *Bioinformatics*, **23**, 365–371.
- Giles, C. and Wren, J. (2008) Large-scale directional relationship extraction and resolution. *BMC Bioinformatics*, **9**, S11.
- Giuliano, C. *et al.* (2006) Exploiting shallow linguistic information for relation extraction from biomedical literature. In *Proceedings of the 11th Conference of the European Chapter of the Association for Computational Linguistics*, Trento.
- Kabiljo, R. *et al.* (2009) A realistic assessment of methods for extracting gene/protein interactions from free text. *BMC Bioinformatics*, **10**, 233.
- Katrenko, S. and Adriaans, P. (2007) Learning Relations from Biomedical Corpora Using Dependency Trees. *Knowl. Discov. Emerg. Compl. Bioinformatics*, **4366**, 61–80.
- Kim, S. *et al.* (2008) Kernel approaches for genic interaction extraction. *Bioinformatics*, **24**, 118–126.
- Kim, S. *et al.* (2010) Walk-weighted subsequence kernels for protein-protein interaction extraction. *BMC Bioinformatics*, **11**, 107.
- Klein, D. and Manning, C.D. (2003) Fast Exact Inference with a Factored Model for Natural Language Parsing. In *Advances in Neural Information Processing Systems 15 (NIPS 2002)*. MIT Press, Cambridge, MA, pp. 3–10.
- Miwa, M. *et al.* (2009a) Protein-protein interaction extraction by leveraging multiple kernels and parsers. *Int. J. Med. Inform.*, **78**, e39–e46.
- Miwa, M. *et al.* (2009b) A rich feature vector for protein-protein extraction from multiple corpora. In *Proceedings of EMNLP*, Singapore, pp. 121–130.
- Miwa, M. *et al.* (2010) Event extraction with complex event classification using rich features. *J. Bioinform. Comput. Biol.*, **8**, 131–146.
- Miyao, Y. *et al.* (2009) Evaluating contributions of natural language parsers to protein-protein interaction extraction. *Bioinformatics*, **25**, 394–400.
- Niu, Y. *et al.* (2009) Evaluation of linguistic features useful in extraction of interactions from PubMed; Application to annotating known, high-throughput and predicted interactions in I2D. *Bioinformatics*, **26**, 111–119.
- Pyysalo, S. *et al.* (2008) Comparative analysis of five protein-protein interaction corpora. *BMC Bioinformatics*, **9**, S6.
- Rinaldi, F. *et al.* (2010) OntoGene in BioCreative II.5. *IEEE/ACM Trans. Comput. Biol. Bioinformatics*, **7**, 472–480.
- Sætre, R. *et al.* (2007) Syntactic features for protein-protein interaction extraction. In *Proceedings of the 2nd International Symposium on Languages in Biology and Medicine*, Singapore.
- Sætre, R. *et al.* (2010) Extracting Protein-Interactions from Text with the Unified AkaneRE Event Extraction System. *IEEE/ACM Trans. Comput. Biol. Bioinform.*, **99**, 442–453.
- Van Landeghem, S. *et al.* (2008) Extracting protein-protein interactions from text using rich feature vectors and feature selection. In *Proceedings of the Third International Symposium on Semantic Mining in Biomedicine (SMBM 2008)*, Turku, Finland, pp. 77–84.
- Witten, I.H. and Frank, E. (2005) *Data Mining: Practical Machine Learning Tools and Techniques*, 2nd edn. Morgan Kaufmann, San Francisco, CA.