

## Fast construction of FM-index for long sequence reads

Heng Li

Medical Population Genetics Program, Broad Institute, 75 Ames Street, Cambridge, MA 02142, USA

Associate Editor: Michael Brudno

## ABSTRACT

**Summary:** We present a new method to incrementally construct the FM-index for both short and long sequence reads, up to the size of a genome. It is the first algorithm that can build the index while implicitly sorting the sequences in the reverse (complement) lexicographical order without a separate sorting step. The implementation is among the fastest for indexing short reads and the only one that practically works for reads of averaged kilobases in length.

**Availability and implementation:** <https://github.com/lh3/ropebwt2>

**Contact:** [hengli@broadinstitute.org](mailto:hengli@broadinstitute.org)

Received on June 2, 2014; revised on July 15, 2014; accepted on August 4, 2014

## 1 INTRODUCTION

FM-index plays an important role in DNA sequence alignment, *de novo* assembly (Simpson and Durbin, 2012) and compression (Cox *et al.*, 2012). Fast and lightweight construction of FM-index for a large dataset is the key to these applications. In this context, a few algorithms (Bauer *et al.*, 2013; Liu *et al.*, 2014) have been developed that substantially outperform earlier algorithms. However, they are only efficient for short reads. A fast and practical algorithm for long sequence reads is still lacking. This work aims to fill this gap.

## 2 METHODS

Let  $\Sigma = \{A, C, G, T, N\}$  be the *alphabet* of DNA with a lexicographical order  $A < C < G < T < N$ . Each element in  $\Sigma$  is called a *symbol* and a sequence of symbols called a *string* over  $\Sigma$ . Given a string  $P$ ,  $|P|$  is its length and  $P[i]$  the symbol at position  $i$ . A sentinel  $\$$  is smaller than all the other symbols. For simplicity, we let  $P[-1] = P[|P|] = \$$ . We also introduce  $\bar{P}$  as the reverse of  $P$  and  $\bar{P}$  as the reverse complement of  $P$ .

Given a list of strings over  $\Sigma$ ,  $(P_i)_{0 \leq i < m}$ , let  $T = P_0 \$ \dots P_{m-1} \$$  with  $\$ < \dots < \$_{m-1} < A < C < G < T < N$ . The *suffix array* of  $T$  is an integer array  $S$  such that  $S(i)$ ,  $0 \leq i < |T|$ , is the starting position of the  $i$ -th smallest suffix in the collection  $T$ . The *Burrows-Wheeler Transform*, or *BWT*, of  $T$  can be computed as  $B[i] = T[S(i) - 1]$ . For the description of the algorithm, we segment  $B$  into  $B = B_\$ B_A B_C B_G B_T B_N$ , where  $B_a[i] = B[i + C(a)]$  with  $C(a) = |\{j : T[j] < a\}|$  being the array of accumulative counts. By the definition of suffix array and BWT,  $B_a$  consists of all the symbols with their next symbol in  $T$  being  $a$ .

The above defines BWT for an ordered list of strings. We next seek to define BWT for an unordered set of strings  $\mathcal{C}$  by imposing an arbitrary sorting order on  $\mathcal{C}$ . We say list  $(P_i)_i$  is in the *reverse lexicographical order* or *RLO*, if  $\bar{P}_i \leq \bar{P}_j$  for any  $i < j$ ; say it is in the *reverse-complement lexicographical order* or *RCLO*, if  $\bar{P}_i \leq \bar{P}_j$  for any  $i < j$ . The *RLO-BWT* of  $\mathcal{C}$ , denoted by  $B^{\text{RLO}}(\mathcal{C})$ , is constructed by sorting strings in  $\mathcal{C}$  in RLO and then applying the procedure in the previous paragraph on the sorted list. *RCLO-BWT*  $B^{\text{RCLO}}(\mathcal{C})$  can be constructed in a similar way.

In  $B^{\text{RCLO}}(\{P_i\}_i \cup \{\bar{P}_j\}_j)$ , the  $k$ -th smallest sequence is the reverse complement of the  $k$ -th sequence in the FM-index. This property removes the necessity of keeping an extra array to link the rank and the position of a sequence in the FM-index, and thus helps to reduce the memory of some FM-index-based algorithms (Simpson and Durbin, 2012). For short reads, RLO/RCLO-BWT is also more compressible (Cox *et al.*, 2012).

As a preparation, we further define two string operations:  $\text{rank}(c, k; B)$  and  $\text{insert}(c, k; B)$ , where  $\text{rank}(c, k; B) = |\{i < k : B[i] = c\}|$  gives the number of symbols  $c$  before the position  $k$  in  $B$ , and  $\text{insert}(c, k; B)$  inserts symbol  $c$  after  $k$  symbols in  $B$  with all the symbols after position  $k$  shifted to make room for  $c$ . We implemented the two operations by representing each  $B_c$  in a  $B^+$ -tree in memory, where a leaf keeps a run-length encoded string and an internal node keeps the count of each symbol in the leaves descended from the node.

Algorithm 1 appends a string to an existing index by inserting each of its symbol from the end of  $P$ . It was first described by Chan *et al.* (2004). Algorithm 2 constructs RLO/RCLO-BWT in a similar manner to Algorithm 1 except that it inserts  $P[i]$  to  $[l, u)$ , the suffix array interval of  $P$ 's suffix starting at  $i + 1$ , and that BWT symbols in this interval are already sorted. This process implicitly applies a radix sort from the end of  $P$ , sorting it into the existing strings in the BWT in RLO/RCLO. Note that if we change line 1 to " $l \leftarrow u \leftarrow |\{i : B[i] = \$\}|$ ", Algorithm 2 will be turned into Algorithm 1. Recall that the BCR algorithm (Bauer *et al.*, 2013) is, to some extent, the multi-string version of Algorithm 1. Following similar reasoning, we can extend Algorithm 2 so as to insert multiple strings at the same time, which gives Algorithm 3. We use an array  $A(j)$  to keep the state of the  $j$ -th sequence after inserting its  $d$ -long suffix. At line 2,  $A(j).c$  is the previously inserted symbol and  $[A(j).l, A(j).u)$  is the interval to which the new symbol is inserted. In implementation, we may speed up the sorting mode by inserting multiple symbols at line 3.

When  $B$  is represented by a balanced tree structure, the time complexity of all three algorithms is  $O(n \log n)$ , where  $n$  is the total number of symbols in the input. However, we will see later that for short strings, Algorithm 3 is substantially faster than the first two algorithms, due to the locality of memory accesses, the possibility of cached  $B^+$ -tree update and the parallelization of the 'for' loop at line 1. These techniques are more effective for a larger batch of shorter strings.

Disregarding RLO/RCLO, Algorithm 3 is similar to BCR except that BCR keeps  $B$  in monolithic arrays. As a result, the time complexity of BCR is  $O(nl)$ , where  $l$  is the maximum length of reads, not scaling well to  $l$ .

## Algorithm 1: Append one string

**Input:** A string  $P$  and an existing BWT  $B$  for  $T$

**Output:** BWT for  $TP\$$

**Function** INSERTIO1( $B, P$ ) **begin**

```

     $c \leftarrow \$$ ;  $k \leftarrow |\{i : B[i] = \$\}|$ 
    for  $i \leftarrow |P| - 1$  to  $-1$  do
         $\text{insert}(P[i], k; B_c)$ 
         $k \leftarrow \text{rank}(P[i], k; B_c) + \sum_{a < c} |\{j : B_a[j] = P[i]\}|$ 
     $c \leftarrow P[i]$ 
    return  $B$ 

```

**Algorithm 2:** Insert one string to RLO/RCLO-BWT

**Input:**  $B^{\text{RLO}}(C)$  (or  $B^{\text{RCLO}}(C)$ ) and a string  $P$   
**Output:**  $B^{\text{RLO}}(C \cup \{P\})$  (or  $B^{\text{RCLO}}(C \cup \{P\})$ )

**Function** INSERTRLO1( $B, P, is\_comp$ ) **begin**  
1  $[l, u] \leftarrow [0, |\{i : B[i] = \$\}|]$   
  **for**  $i \leftarrow |P| - 1$  **to**  $-1$  **do**  
     $[l, u] \leftarrow \text{INSERTAUX}(B, P[i], l, u, P[i + 1], is\_comp)$   
  **return**  $B$

**Function** INSERTAUX( $B, c', l, u, c, is\_comp$ ) **begin**  
 $k \leftarrow l$   
  **if**  $is\_comp$  **is true and**  $c' \neq \text{"N"}$  **then**  
    **for**  $a = \$$  **or**  $c' < a < \text{"N"}$  **do**  
       $k \leftarrow k + [\text{rank}(a, u; B_c) - \text{rank}(a, l; B_c)]$   
  **else**  
    **for**  $\$ \leq a < c'$  **do**  
       $k \leftarrow k + [\text{rank}(a, u; B_c) - \text{rank}(a, l; B_c)]$   
   $l' \leftarrow \text{rank}(c', l; B_c); u' \leftarrow \text{rank}(c', u; B_c)$   
   $\text{insert}(c', k; B_c)$   
   $m \leftarrow \sum_{a < c} |\{j : B_a[j] = c'\}|$   
  **return**  $[l' + m, u' + m]$

**Algorithm 3:** Insert multiple strings

**Input:** Existing BWT  $B$  and a list of strings  $\{P_k\}_k$   
**Output:** Updated BWT  $B$  with strings inserted in the specified order

**Function** INSERTMULTI( $B, \{P_k\}_k, is\_sorted, is\_comp$ ) **begin**  
  **for**  $0 \leq j < |\{P_k\}_k|$  **do**  
     $A(j).c \leftarrow \$; A(j).i \leftarrow j$   
    **if**  $is\_sorted$  **is true then**  
       $[A(j).l, A(j).u] \leftarrow [0, |\{i : B[i] = \$\}|]$   
    **else**  
       $A(j).l \leftarrow A(j).u \leftarrow |\{i : B[i] = \$\}| + j$   
   $d \leftarrow 0$   
  **while**  $|A| \neq 0$  **do**  
1    $\text{Stable sort array } A \text{ by } A(\cdot).c$   
2   **for**  $0 \leq j < |A|$  **do**  
     $c \leftarrow A(j).c; A(j).c \leftarrow P_{A(j).i}[|P_{A(j).i}| - 1 - d]$   
     $[A(j).l, A(j).u]$   
3    $\leftarrow \text{INSERTAUX}(B, A(j).c, A(j).l, A(j).u, c, is\_comp)$   
     $\text{Remove } A(j) \text{ if } A(j).c = \$$   
     $d \leftarrow d + 1$   
  **return**  $B$

### 3 RESULTS AND DISCUSSION

We implemented the algorithm in ropeBWT2 and evaluated its performance together with BEETL (<http://bit.ly/beetlGH>), the original on-disk implementation of BCR and BCRext, ropeBWT-BCR (<https://github.com/lh3/ropebwt>), an in-memory reimplementation of BCR by us, and NVBio (<http://bit.ly/nvbioio>), a GPU-based algorithm inspired by CX1 (Liu *et al.*, 2014). Table 1 shows that for  $\sim 100$  bp reads, ropeBWT2 has comparable performance to others. For the  $\sim 875$  bp Venter dataset, NVBio aborted due to insufficient memory under various settings. We did not apply BCR because it is not designed for

**Table 1.** Performance of BWT construction

| Data <sup>a</sup> | Algorithm    | RCLO | Real   | CPU% | RAM <sup>b</sup><br>(GB) | Comments              |
|-------------------|--------------|------|--------|------|--------------------------|-----------------------|
| worm              | nvbio        | —    | 316 s  | 138  | 12.9                     | See note <sup>c</sup> |
| worm              | ropebwt-bcr  | —    | 480 s  | 223  | 2.2                      | -btORf                |
| worm              | Algorithm 3  | Yes  | 506 s  | 250  | 10.5                     | -brRm10g              |
| worm              | Algorithm 3  | No   | 647 s  | 249  | 11.8                     | -bRm10g               |
| worm              | beetl-bcr    | —    | 965 s  | 259  | 1.8                      | RAM disk <sup>d</sup> |
| worm              | beetl-bcr    | —    | 2092 s | 122  | 1.8                      | Network <sup>e</sup>  |
| worm              | Algorithm 1  | —    | 5125 s | 100  | 2.5                      | -bRm0                 |
| worm              | beetl-bcrext | —    | 5900 s | 48   | 0.1                      | Network <sup>e</sup>  |
| 12 878            | ropebwt-bcr  | —    | 3.3 h  | 210  | 39.3                     | -btORf                |
| 12 878            | nvbio        | —    | 4.1 h  | 471  | 63.8                     | See note <sup>f</sup> |
| 12 878            | Algorithm 3  | Yes  | 5.0 h  | 261  | 34.0                     | -brRm10g              |
| 12 878            | Algorithm 3  | No   | 5.1 h  | 248  | 60.9                     | -bRm10g               |
| 12 878            | beetl-bcr    | —    | 11.2 h | 131  | 31.6                     | Network <sup>e</sup>  |
| Venter            | Algorithm 3  | Yes  | 1.4 h  | 274  | 22.2                     | -brRm10g              |
| Venter            | Algorithm 3  | No   | 1.5 h  | 274  | 22.8                     | -bRm10g               |
| mol               | Algorithm 3  | No   | 6.8 h  | 285  | 20.0                     | -bRm10g               |

<sup>a</sup>Datasets—worm: 66M  $\times$  100 bp *Caenorhabditis elegans* reads from SRR065390; 12878: 1206M  $\times$  101 bp human reads for sample NA12878 (Depristo *et al.*, 2011). Venter: 32M  $\times$  875 bp (in average) human reads by Sanger sequencing (Levy *et al.* 2007; <http://bit.ly/levy2007>); mol: 23M  $\times$  4026 bp (in average) human reads by Illumina's Molecule sequencing (<http://bit.ly/mol12878>).

<sup>b</sup>Hardware—CPU: 48 cores of Xeon E5-2697v2 at 2.70 GHz; GPU: one Nvidia Tesla K40; RAM: 128 GB; Storage: Isilon IQ 72000x and X400 over network. CPU time, wall-clock time and peak memory are measured by GNU time.

<sup>c</sup>Run with option '-R -cpu-mem 4096 -gpu-mem 4096'. NVBio uses more CPU and GPU RAM than the specified.

<sup>d</sup>Results and temporary files created on in-RAM virtual disk '/dev/shm'.

<sup>e</sup>Results and temporary files created on Isilon's network file system.

<sup>f</sup>Run with option '-R -cpu-mem 48000 -gpu-mem 4096'.

long reads of unequal lengths. Only ropeBWT2 works with this data set and the even longer molecule reads.

**Funding:** NHGRI U54HG003037; NIH GM100233.

**Conflict of Interest:** none declared.

### REFERENCES

- Bauer, M.J. *et al.* (2013) Lightweight algorithms for constructing and inverting the BWT of string collections. *Theor. Comput. Sci.*, **483**, 134–148.
- Chan, H.-L. *et al.* (2004) Compressed index for a dynamic collection of texts. In: Sahinalp, S.C. Muthukrishnan, S. and Dogrusöz, U. (eds) *CPM, Volume 3109 of Lecture Notes in Computer Science*. Springer, Berlin Heidelberg, pp. 445–456.
- Cox, A.J. *et al.* (2012) Large-scale compression of genomic sequence databases with the burrows-wheeler transform. *Bioinformatics*, **28**, 1415–1419.
- Depristo, M.A. *et al.* (2011) A framework for variation discovery and genotyping using next-generation DNA sequencing data. *Nat. Genet.*, **43**, 491–498.
- Levy, S. *et al.* (2007) The diploid genome sequence of an individual human. *PLoS Biol.*, **5**, e254.
- Liu, C.-M. *et al.* (2014) GPU-accelerated BWT construction for large collection of short reads. arXiv:1401.7457.
- Simpson, J.T. and Durbin, R. (2012) Efficient de novo assembly of large genomes using compressed data structures. *Genome Res.*, **22**, 549–556.