GPU accelerated biochemical network simulation

Yanxiang Zhou[1,†], Juliane Liepe[1], Xia Sheng[1,2], Michael P. H. Stumpf[1,2,3,*]
and Chris Barnes[1,2,*]

[1]Centre for Bioinformatics, Division of Molecular Biosciences, [2]Institute of Mathematical Sciences and [3]Centre for
Integrative Systems Biology, Imperial College London, London, UK

Associate Editor: Olga Troyanskaya

ABSTRACT

Motivation: Mathematical modelling is central to systems and
synthetic biology. Using simulations to calculate statistics or to
explore parameter space is a common means for analysing these
models and can be computationally intensive. However, in many
cases, the simulations are easily parallelizable. Graphics processing
units (GPUs) are capable of efficiently running highly parallel
programs and outperform CPUs in terms of raw computing power.
Despite their computational advantages, their adoption by the
systems biology community is relatively slow, since differences in
hardware architecture between GPUs and CPUs complicate the
porting of existing code.
Results: We present a Python package, *cuda-sim*, that provides
highly parallelized algorithms for the repeated simulation of
biochemical network models on NVIDIA CUDA GPUs. Algorithms
are implemented for the three popular types of model formalisms: the
LSODA algorithm for ODE integration, the Euler–Maruyama algorithm
for SDE simulation and the Gillespie algorithm for MJP simulation.
No knowledge of GPU computing is required from the user. Models
can be specified in SBML format or provided as CUDA code. For
running a large number of simulations in parallel, up to 360-fold
decrease in simulation runtime is attained when compared to single
CPU implementations.
Availability: http://cuda-sim.sourceforge.net/
Contact: christopher.barnes@imperial.ac.uk; m.stumpf@imperial.ac.uk
Supplementary information: Supplementary data are available at
Bioinformatics online.

1 INTRODUCTION

Mathematical modelling is an integral part of systems and synthetic
biology. Ordinary differential equations (ODEs) are the most
commonly used methodology, but due to increasing appreciation
of the importance of stochasticity in biological processes, stochastic
differential equations (SDEs) and Markov jump processes (MJPs)
are also applied. Since most models are non-linear, they generally
cannot be solved analytically and therefore require numerical
treatment. Furthermore, in order to understand behaviour across
high-dimensional parameter space or to perform simulation based
inference (Liepe *et al.*, 2010; Toni *et al.*, 2009), a very large number

---

*To whom correspondence should be addressed.

†Present address: Apoptosis and Proliferation Control Laboratory, Cancer
Research UK, London Research Institute, London, UK.

of simulations is required, making analysis of even the simplest
models extremely time consuming.

For computationally expensive calculations, graphics processing
units (GPUs) can be used. GPUs are many core, multi-threaded
chips that are capable of several hundred GFLOPS (Kirk and
Hwu, 2010). In terms of raw computing power, a single GPU in
a desktop PC is comparable to a CPU cluster, but is much cheaper.
However, due to their single instruction multiple data (SIMD)
architecture, only highly parallelized processes can be efficiently
run on GPUs. Even though platforms for general purpose GPU
computing like CUDA® (Compute Unified Device Architecture)
from NVIDIA® which provides a CUDA API exist, it remains
difficult and time consuming to port existing algorithms that were
designed for execution on CPUs. There have been developments
in porting biochemical network simulators to GPUs (reviewed in
Dematté and Prandi, 2010) but there does not currently exist a
general purpose simulation tool that integrates multiple algorithms
within the same interface.

Here, we present a new Python package called *cuda-sim* which
provides highly parallelized algorithms for large scale simulations
of biochemical network models. It is compatible with all NVIDIA
GPUs that support CUDA. Absolutely no knowledge of CUDA and
GPU computing is needed for the user to access the simulation
algorithms: (1) The integration of ODEs is carried out using a GPU
implementation of LSODA (Hindmarsh, 1983), (2) SDE simulations
are provided via the Euler-Maruyama algorithm (Kloeden and
Platen, 1999) and (3) simulations from a MJP (or Master equation)
are performed using the Gillespie algorithm (Gillespie, 1976). All
functionality can be accessed via a Python interface that hides the
implementation details.

2 IMPLEMENTATION

The *cuda-sim* package is implemented in Python using PyCUDA
(Klöckner *et al.*, 2009) to access the GPU. PyCUDA acts as a
wrapper for the CUDA API and provides abstractions that facilitate
programming. As an additional layer between code and hardware,
PyCUDA can also accommodate future changes in GPU architecture
and will help ensuring that *cuda-sim* remains compatible with newer
GPU generations. The package is written in an object oriented
manner with an abstract simulator base class such that there is a
common interface for accessing the algorithms.

Two different pseudo random number generators (RNG) are used
for the SDE and MJP simulations. In the MJP simulations, each
thread carries out different numbers of simulation steps and therefore
needs different numbers of random numbers that are provided by
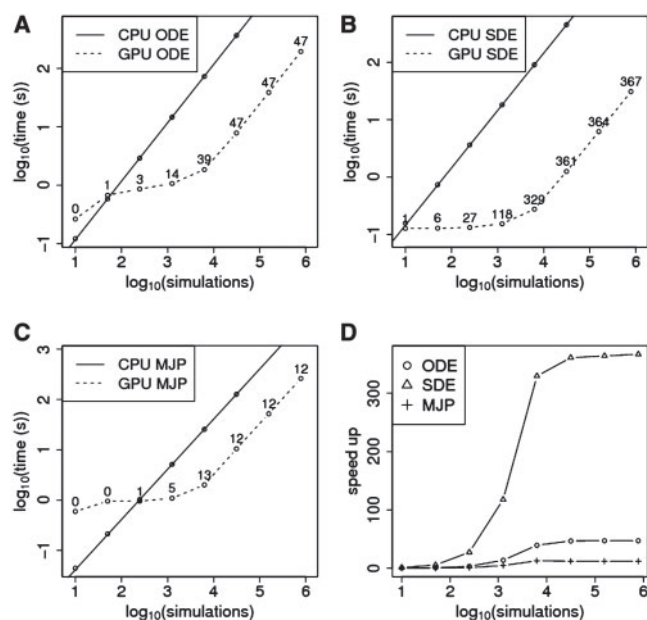
Fig. 1. Timing comparisons. (**A–C**) Time taken to simulate a given number of realisations for a single core of an Intel Core i7-975 Extreme Edition Processor 3.33 GHz (solid line) and one Tesla C2050 GPU (dashed line) for (A) the LSODA (B) the Euler–Maruyama and (C) the Gillespie algorithm, respectively. The relative speed-ups for given numbers of simulations are indicated next to the GPU simulation results. (**D**) Summary of the relative speed-up of the three different algorithms.

one Mersenne Twister RNG (Matsumoto and Nishimura, 1998) per thread. In the completely parallel SDE simulations, each thread requires the same number of random numbers at the same time. This fact is exploited by a binary linear operations RNG that is local to a group of threads known as a *warp* (Thomas *et al.*, 2009).

The user accesses the package by specifying SBML models which are parsed in *cuda-sim* using libSBML (Bornstein *et al.*, 2008) and then converted to CUDA code modules. These code modules are automatically incorporated into the GPU kernels and run from within *cuda-sim*. Using a provided Python script, the user can specify model parameters and run the simulations with *cuda-sim*. Alternatively, the algorithms can be called directly from Python, allowing incorporation into other software projects.

It is advisable to use dedicated general purpose GPUs like the Tesla C2050 that we used for our timing studies. GPUs that provide graphical output have a time limitation on the execution length of programs and therefore will not be compatible with larger simulations run with *cuda-sim*.

3 TIMING COMPARISONS

Using a model of p53-Mdm2 oscillations which contains 3 species, 6 reactions and 8 parameters (Geva-Zatorsky *et al.*, 2006 and Supplementary Fig. S1) we simulated different numbers of time series of length 100 h and performed timing studies comparing the runtime on the GPU and the runtime on a *single* CPU (Fig. 1). For the ODE integrations to yield different results for each thread, instead of

using fixed parameters, we drew parameters from uniform random distributions in the interval between 0 and 2.

The CPU versions of the Gillespie and Euler–Maruyama algorithms were written in C++ and compiled with GCC using the -O3 optimization flag. For the LSODA algorithm comparisons, the CPU implementation in the SciPy Python module was used. The testing was done on an Intel Core i7-975 Extreme Edition Processor 3.33 GHz machine with 12 GB of RAM and one Tesla C2050 GPU. All points are averages over three runs. Since the CPU runtimes scale linearly, the total CPU time for large numbers of simulations can be extrapolated using a linear model.

For the LSODA, the Euler–Maruyama and the Gillespie algorithm, speed-ups of 47-fold, 367-fold and 12-fold are attained, respectively, for large numbers of simulations. Only for small numbers of simulations, are the CPU implementations of the three algorithms faster than the GPU versions (Fig. 1 A–C). This is due to the fact that the initialization on the GPU takes substantially longer than on the CPU. But since in most applications of these algorithms, either in order to explore the parameter space or to perform inference, at least thousands of simulations will be needed for which the GPU outperforms the CPU even for the rather simple p53-Mdm2 model.

We also compared the *cuda-sim* implementations of the LSODA and Gillespie algorithms with implementations in the Matlab package *SBTOOLBOX2* (Schmidt and Jirstand, 2006) and our Euler–Maruyama implementation with the native *sde* function within Matlab. Since stochastic simulation in *SBTOOLBOX2* supports only mass-action models, we used a model of enzyme kinetics (Supplementary Fig. S2). We obtained similar timing to the p53-Mdm2 model when using our CPU implementations (Supplementary Fig. S3) and speed-ups of between three and four orders of magnitude when compared to Matlab (Supplementary Fig. S4).

4 CONCLUSIONS

GPUs offer a powerful and cost-effective solution for parallel computing. *cuda-sim* provides a Python interface for biochemical network simulations using ODEs, SDEs and MJPs on NVIDIA CUDA GPUs and significantly reduces computation time. The package can be used as a standalone tool, or incorporated into other Python packages.

Conflict of Interest: none declared.

REFERENCES

Bornstein,B.J. *et al.* (2008) LibSBML: an API Library for SBML. *Bioinformatics*, **24**, 880–881.

Dematté,L. and Prandi,D. (2010) GPU computing for systems biology. *Brief Bioinform.*, **11**, 323–333.

Geva-Zatorsky,N. *et al.* (2006) Oscillations and variability in the p53 system. *Mol. Syst. Biol.*, **2**, 2006.0033.

Gillespie,D.T. (1976) A general method for numerically simulating the stochastic time evolution of coupled chemical reactions. *J. Comput. Phy.*, **22**, 403–434.

Hindmarsh,A.C. (1983) ODEPACK, A Systematized Collection of ODE Solvers, Scientific Computing. In Stepleman,R.S. *et al*. (eds) *IMACS Transactions on Scientific Computation*. Vol. 1. Amsterdam, North-Holland, pp. 55-64.

Kirk,D.B. and Hwu,W.W. (2010) *Programming Massively Parallel Processors*. Morgan Kaufmann, Burlington.

Klöckner,A. *et al*. (2009) PyCUDA: GPU run-time code generation for high-performance computing. arXiv:0911.3456

Kloeden,P.E. and Platen,E. (1999) *Numerical Solution of Stochastic Differential Equations*. Springer, Berlin and Heidelberg.

Liepe,J. *et al*. (2010) ABC-SysBio–approximate Bayesian computation in Python with GPU support. *Bioinformatics*, **26**, 1797–1799.

Matsumoto,M. and Nishimura,T. (1998) Mersenne twister: a 623-dimensionally equidistributed uniform pseudo-random number generator. *ACM Trans. Model. Comput. Simul.*, **8**, 3–30.

Schmidt,H. and Jirstand,M. (2006) Systems Biology Toolbox for MATLAB: a computational platform for research in systems biology. *Bioinformatics*, **22**, 514–515.

Thomas,D. *et al*. (2009) A comparison of CPUs, GPUs, FPGAs and massively parallel processor arrays for random number generation. In *Proceeding of the ACM/SIGDA international symposium on Field programmable gate arrays*. ACM, New York.

Toni,T. *et al*. (2009) Approximate Bayesian computation scheme for parameter inference and model selection in dynamical systems. *J. R. Soc. Interface*, **6**, 187–202.