mRMRe: an R package for parallelized mRMR ensemble feature selection

Nicolas De Jay[1,†], Simon Papillon-Cavanagh[1,†], Catharina Olsen[2], Nehme El-Hachem[1], Gianluca Bontempi[2] and Benjamin Haibe-Kains[1,*]

[1]Bioinformatics and Computational Biology Laboratory, Integrative Systems Biology Axis, Institut de recherches cliniques de Montréal, Montreal, H2W 1R7, Quebec, Canada and [2]Machine Learning Group, Department of Computer Science, Université Libre de Bruxelles, 1050, Brussels, Belgium

Associate Editor: Olga Troyanskaya

ABSTRACT

Motivation: Feature selection is one of the main challenges in analyzing high-throughput genomic data. Minimum redundancy maximum relevance (mRMR) is a particularly fast feature selection method for finding a set of both relevant and complementary features. Here we describe the mRMRe R package, in which the mRMR technique is extended by using an ensemble approach to better explore the feature space and build more robust predictors. To deal with the computational complexity of the ensemble approach, the main functions of the package are implemented and parallelized in C using the openMP Application Programming Interface.

Results: Our ensemble mRMR implementations outperform the classical mRMR approach in terms of prediction accuracy. They identify genes more relevant to the biological context and may lead to richer biological interpretations. The parallelized functions included in the package show significant gains in terms of run-time speed when compared with previously released packages.

Availability: The R package mRMRe is available on Comprehensive R Archive Network and is provided open source under the Artistic-2.0 License. The code used to generate all the results reported in this application note is available from Supplementary File 1.

Contact: bhaibeka@ircm.qc.ca

Supplementary information: Supplementary data are available at *Bioinformatics* online.

1 INTRODUCTION

In genomic data analysis, phenotype-associated feature selection is of utmost importance in understanding the biological processes underlying the relevant phenotype and in building accurate predictive models. This is not a trivial task for high-throughput genomic data (microarray, next-generation sequencing, etc.), as they are high dimensional, noisy and have complex intercorrelational structures. Because feature selection is an nondeterministic polynomial time (NP)-hard problem, one must resort to the use of heuristics, which find suitable and sub-optimal sets of relevant features in high-dimensional datasets. Among these heuristics,

the minimum redundancy maximum relevance (mRMR) feature selection technique is particularly appealing because of the relatively low computational complexity of its algorithm for finding a set of relevant and complementary features (Ding and Peng, 2005), from which accurate predictive models are developed. The challenge lies in that mRMR, like all feature selection algorithms in a low sample-to-dimensionality ratio setting, produces highly variable results, and small changes in sample data often lead to dramatically different sets of selected features.

In the past decade, a new trend has emerged whereby highly accurate 'ensemble' classifiers are produced by combining less accurate ones, on the grounds that model variance is reduced without affecting bias (Kittler *et al.*, 1998). However, the computational cost of an ensemble variant of the mRMR method is high, as multiple mRMR feature selections must be done. This was the primary motivator for the development of a new R package, mRMRe, which implements an ensemble variant of mRMR, in which multiple feature sets, rather than a single list of features, is built. Also included in the package is a function for computing a mutual information matrix (MIM) based on the appropriate estimators for each variable type (i.e. continuous, discrete and survival data). Both these package features have been adapted to fully use multicore platforms.

2 METHODS

The *mim* function computes a MIM using a linear approximation based on correlation such that mutual information (MI) is estimated as $I(x, y) = -\frac{1}{2}\ln(1 - \rho(x, y)^2)$, where I and ρ, respectively, represent the MI and correlation coefficient between variables x and y. Correlation between continuous variables can be computed using either Pearson's or Spearman's estimators, whereas Cramer's V is used for correlation between discrete variables and Somers' Dxy index is used for correlation between continuous variables and survival data.

The mRMR technique, as implemented in the *mRMR.classic* function, allows an efficient selection of relevant and non-redundant features (Ding and Peng, 2005). Let y be the output variable and $X = \{x_1, \ldots, x_n\}$ be the set of n input features. The method ranks X by maximizing the MI with y (maximum relevance) and minimizing the average MI with all the previously selected variables (minimum redundancy). Let x_i be the feature with highest MI with the output variable and thus selected first

$$x_i = \arg \max_{x_i \in X} I(x_i, y) \tag{1}$$

---

*To whom correspondence should be addressed.
†The authors wish it to be known that, in their opinion, the first two authors should be regarded as joint First Authors.

Table 1. Run-time performances in seconds of mRMRe and *sideChannelAttack* packages using the CGP dataset

Package	Full MIM	mRMR	mRMRe.e	mRMRe.b
mRMRe	80	23	36	863
sideChannelAttack	158	212	NA	NA

Note: Benchmarks were performed on a Dell PowerEdge R815 using 8 CPU cores (out of 48, from 4×12-core, 2.2 GHz AMD Opteron) and 12 GB of RAM.

The set of selected features, denoted by S, is then initialized with x_i. Next, another feature is added to S by choosing the feature having the highest relevance with the output variable and the lowest redundancy with the previously selected features, thus maximizing the score q at step j

$$q_j = I(x_j, y) - \frac{1}{|S|} \sum_{x_k \in S} I(x_j, x_k) \qquad (2)$$

This step is repeated until the desired solution length has been attainted. We have implemented this approach for continuous/survival and discrete variables, also referred to as F-test Correlation Difference (FCD) and Mutual Information Difference (MID) schemes in Table 1 of (Ding and Peng, 2005), respectively.

Although mRMR is a fast and greedy heuristic, it is not guaranteed to find a global optimal solution should one exist. Alternative feature subsets of equivalent or better quality than the one identified may exist. Moreover, the features selected by a single mRMR run are unlikely to adequately account for the diversity of the biological processes associated with the phenotype under study.

To alleviate these problems we implemented two ensemble approaches to generate multiple mRMR solutions in parallel; these two techniques are referred to as *exhaustive* and *bootstrap* ensemble mRMR. The exhaustive variant extends the classical mRMR heuristic by initializing multiple feature selection procedures with the $k > 1$ most relevant features. Subsequently, k mRMR solutions are produced in parallel, in which the first selected feature is guaranteed to be different. The pseudocode describing the algorithm implementing the exhaustive ensemble mRMR feature selection is provided in Algorithm 1 in Supplementary Material.

The bootstrap variant resamples (with replacement) the original dataset to generate k bootstraps, and classical mRMR feature selection is performed in parallel for each of the bootstrapped datasets, thus generating k mRMR solutions. The pseudocode describing the algorithm implementing the bootstrap ensemble mRMR feature selection is provided in Algorithm 2 in Supplementary Material.

A considerable share of the computational complexity of existing mRMR packages is due to inefficient computation of the MIM. These, such as *minet* (Meyer *et al.*, 2008) and *sideChannelAttack* (Lerman *et al.*, 2011), compute the MIM completely before performing the mRMR feature selection. However, only a small portion of the MIM is generally required to compute mRMR scores (Equations 1 and 2) during the feature selection process. We accelerated our mRMR implementation by computing the MI score between features in a lazy-evaluation manner, computing them only when needed. This significantly reduces the run-time complexity of mRMR and is a critical point for controlling the computational time of the ensemble variants.

3 CASE STUDY

In this case study, we assess the benefits of using the exhaustive and bootstrap ensemble mRMR feature selection methods (referred to as *mRMRe.e* and *mRMRe.b*, respectively) by analyzing two recently published pharmacogenomic datasets generated

within the Cancer Genome Project [CGP; (Garnett *et al.*, 2012)] and the Cancer Cell Lines Encyclopedia [CCLE; (Barretina *et al.*, 2012)]. In these large datasets of cancer cell lines, the authors measured sensitivity (IC_{50}) to Irinotecan (Camptothecin), a therapeutic drug mainly used in colon cancer. This metric was used to discriminate between resistant and sensitive cell lines. Similar to Papillon-Cavanagh *et al.* (2013), we used CGP as a training set, whereas CCLE was split in two validation sets: CCLE COMMON contains cell lines common to both CGP and CCLE (471), whereas CCLE NEW contains cell lines absent in CGP (565).

We then implemented five feature selection methods: (i) SINGLEGENE and (ii) RANK consist in ranking the features by correlation with drug sensitivity so as to, respectively, select the first and the top n features; (iii) mRMR is used to select the most relevant, less redundant set of n features associated to drug sensitivity; (iv) mRMRe.e and (v) mRMRe.b perform multiple mRMR feature selections in parallel to identify 200 mRMR solutions using the exhaustive and bootstrap approach, respectively. Feature selection was followed by linear regression model fitting using the selected features to predict drug sensitivity. For mRMRe methods, drug sensitivity is predicted by averaging predictions of the 200 multivariate models corresponding to each mRMR solution.

To strike a balance between model complexity, considered here as the number of selected features (solution length), and prediction performance, we performed 100 resamplings of the training set and estimated the concordance index (Harrell *et al.*, 1996) of the resulting predictive model with respect to feature selection method and solution length (Fig. 1A). As the concordance index is a generalization of the area under the receiver-operating characteristic curve, high-performing models are associated with index values close to 1, and random models are expected to yield index values close to 0.5. We observed that mRMRe methods yielded higher performance (Wilcoxon signed rank test $P < 0.001$, see Supplementary Table S1), especially for small (≤ 5) solution lengths. As expected, the gain in performance rapidly diminished with increasing solution length (Supplementary Fig. S1). We consequently selected 15 as the solution length, exhibiting the most balance between model complexity and performance (Fig. 1B). In addition, we computed the variance of the concordance index over multiple resamplings (Fig. 1B and Supplementary Fig. S2) and observed that the classical and ensemble mRMR variants produced lower variance when compared with RANK or SINGLEGENE. Variance was much lower for ensemble mRMR techniques when compared with classical mRMR techniques for small (≤ 5) solution lengths. However, no difference was observed for larger solution lengths (Supplementary Fig. S2).

In addition to performance assessment diversity (Tsymbal *et al.*, 2005) and stability (Guzmán-Martínez and Alaiz-Rodríguez, 2011; Kuncheva, 2007) of the mRMRe, feature selection techniques implemented were investigated (Supplementary Fig. S3). Although RANK and mRMR select 30 different genes during each resampling of the training set, the mRMRe techniques identify multiple, possibly diverse, mRMR solutions in parallel. In fact, we observed that, on average, mRMRe.e selects 210 distinct features shared between the 200 distinct mRMR solutions (Supplementary Fig. S3A), whereas
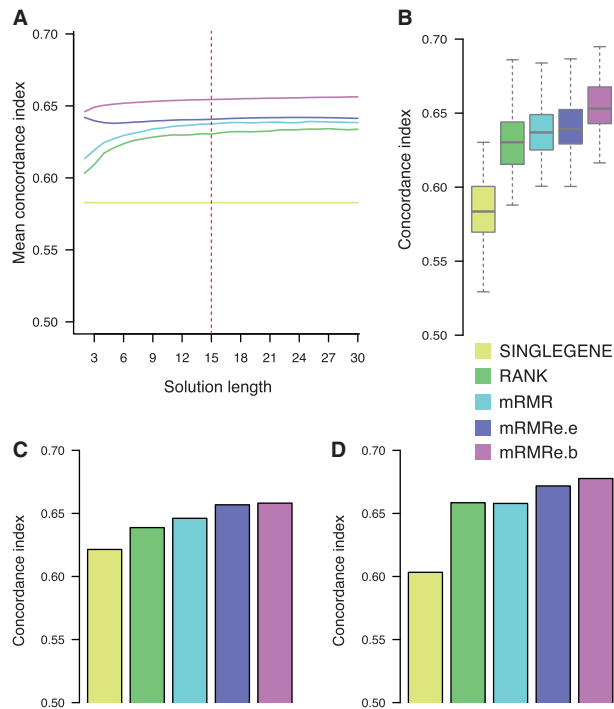
Fig. 1. Performance of the five feature selection techniques. (**A**) Mean concordance index with respect to feature selection technique, and solution length in 100 resamplings of the training set (CGP). The dashed vertical bar in red is the chosen cutoff between solution length and model performance, the 15 feature mark. (**B**) Box plot representing the concordance indices during resamplings when 15 features are selected. (**C**) Concordance index in CCLE COMMON, validation set composed of the cell lines also present in the training set. (**D**) Concordance index in CCLE NEW, validation set composed of the cell lines not present in the training set

mRMRe.b is much more diverse with an average of 988 distinct features (Supplementary Fig. S3B). The greater diversity of mRMRe.b can partly explain the better performance observed in the training set (Tsymbal *et al.*, 2005). To assess the stability of the five feature selection methods, the Tanimoto Stability Index was implemented as described in (Guzmán-Martínez and Alaiz-Rodríguez, 2011). This index enables the comparison of the sets of selected features with each resampling of the training set. As can be seen in Supplementary Figure S3C, selecting the gene that correlates most with Irinotecan sensitivity yielded the most stable results (Supplementary File S2). Concurring with Garnett *et al.* (2012), we found the most frequently selected gene to be FLI1. However, the association between FLI1 expression and Irinotecan sensitivity remains poorly understood. Ensemble mRMR feature selections produce more stable gene lists when compared with classical mRMR. Stability of the gene lists obtained with mRMRe.e is close to that of SINGLEGENE, whereas mRMRe.b's exhibited stability is virtually identical to that of RANK for solution lengths >7.

To validate the performance of the mRMRe feature selection approach, independent validation sets were tested, and it was observed that mRMRe yielded a higher concordance index overall (Fig. 1C and D). When predicting response to Irinotecan for the same cell lines (CCLE COMMON), mRMR.e and mRMR.b

approaches outperformed classical mRMR with an improvement of 1.7% and 1.9%, respectively (Fig. 1C). For new cell lines (CCLE NEW), the gain was more notable with mRMRe.e and mRMRe.b, yielding an improvement of 2.1% and 3.4%, respectively (Fig. 1D). This supports the superior generalization ability of the ensemble mRMR approaches, especially mRMRe.b because it identifies a more diverse panel of mRMR solutions (Supplementary Fig. S4).

We further investigated the advantage of using mRMRe for biological interpretation by performing a Gene Ontology analysis using the *GOSim* package. Because ensemble mRMR techniques, by their very nature, select more genes than the other techniques, we expect more biological terms to be significantly enriched with the former. This is indeed the case, as illustrated in Supplementary Figure S5: 100, 21, 167 and 400 Gene Ontology terms have been identified as significantly enriched (Fisher's exact test $P < 0.05$) for RANK, mRMR, mRMRe.e and mRMRe.b, respectively (Supplementary File 3).

Finally, we compared the run-time performance of our functions with those of the *sideChannelAttack* package (Lerman *et al.*, 2011); the latter implements the classical mRMR and relies on the *minet* package (Meyer *et al.*, 2008) to build the MIM. As reported in Table 1, our package, on a 8-core workstation, is twice as fast for full MIM construction and 9.2 times faster for classical mRMR feature selection; this is mainly because of the lazy MIM implementation. It is worth noting that the search for 200 mRMR solutions using mRMRe.e is still 5.9 times faster than the classical mRMR with *sideChannelAttack*. Moreover, performing ensemble mRMR feature selection using the bootstrap method is as computationally demanding, as a new (lazy) MIM must be computed for each bootstrap.

4 CONCLUSION

The R package mRMRe provides functions to efficiently perform ensemble mRMR feature selection by taking full advantage of parallel computing. Ensemble mRMR can be beneficial from both a predictive (lower bias and lower variance) and biological (more thorough feature space exploration) point of view, which makes it particularly attractive for high-throughput genomic data analysis.

Conflict of Interest: none declared.

REFERENCES

Barretina,J. *et al.* (2012) The Cancer Cell Line Encyclopedia enables predictive modelling of anticancer drug sensitivity. *Nature*, **483**, 603–607.

Ding,C. and Peng,H. (2005) Minimum redundancy feature selection from microarray gene expression data. *J. Bioinform. Comput. Biol.*, **3**, 185–205.

Garnett,M.J. *et al.* (2012) Systematic identification of genomic markers of drug sensitivity in cancer cells. *Nature*, **483**, 570–575.

Guzmán-Martínez,R. and Alaiz-Rodríguez,R. (2011) Feature selection stability assessment based on the Jensen-Shannon divergence. In: *ECML PKDD 2011*. Springer, Berlin, Heidelberg, pp. 597–612.

Harrell,F.E. Jr *et al.* (1996) Multivariable prognostic models: issues in developing models, evaluating assumptions and adequacy, and measuring and reducing errors. *Stat. Med.*, **15**, 361–387.

Kittler,J. *et al.* (1998) On combining classifiers. *IEEE Trans. Pattern Anal. Mach. Intell.*, **20**, 226–238.

Kuncheva,L.I. (2007) A stability index for feature selection. In: *AIAP'07: Proceedings of the 25th conference on Proceedings of the 25th IASTED International Multi-Conference*. pp. 390–395, Anaheim, CA, USA.

Lerman,L. *et al.* (2011) Side channel attack: an approach based on machine learning. In: *COSADE 2011, second International Workshop on Consecutive Side-Channel Analysis and Secure Design*. Center for Advanced Security Research Darmstadt, Darmstadt, Germany.

Meyer,P.F. *et al.* (2008) minet: a R/Bioconductor package for inferring large transcriptional networks using mutual information. *BMC Bioinformatics*, **9**, 461.

Papillon-Cavanagh,S. *et al.* (2013) Comparison and validation of genomic predictors for anticancer drug sensitivity. *J. Am. Med. Inform. Assoc.*, **20**, 597–602.

Tsymbal,A. *et al.* (2005) Diversity in search strategies for ensemble feature selection. *Inf. Fusion*, **6**, 83–98.